

How Nexora AI actually works

A spreadsheet with five fixed tabs became a system that generates as many levels as a role needs. This is where the data goes, what the machine writes, and the three places it has to stop and wait for a person.


LINEAGE	Interview Assessment Sheet v3.1
SURFACES	6 screens · 2 sub-screens
ROLES	4, gated at the route
MODULES	33 · ~11,850 lines
MUTATIONS	26 reducer actions
HUMAN GATES	3, none bypassable

THE SHAPE OF IT

Four planes, one direction of travel

Everything in the platform sits in one of four planes, and dependencies only ever point downward. A screen can read the record; the record has never heard of a screen.


That constraint is what makes the interesting part possible. Because the engine knows nothing about the surfaces, adding a fifth interview level is a data change, not a new screen.



FLOW 01 – THE SPINE

Application to offer, end to end

This is the path every candidate takes. Five of the six stages are automatic. The three amber bars are the places the system deliberately refuses to continue on its own.




■ Generated per level
 ■ Machine-authored
 ■ Blocking human gate

FLOW 02 – THE GENERATIVE CORE

One declaration, N levels

This is the single architectural decision the whole product turns on. The requisition declares its workflow as data. Interviews, scorecards and the consolidated rollup are all generated from that declaration — nothing downstream is written by hand.

The assessment sheet we started from has five fixed tabs: TA Partner, Business 1, Business 2, HR Partner, Segment Head. A sixth round meant a new tab and a new set of formulas. Here the same code path produces four levels for the Bangalore sales role and five for the Singapore finance role, because those requisitions declared different arrays.




WHAT THIS BUYS THE TEAM

A hiring manager who wants a fifth round adds a stage from the workflow tab. The matching scorecard appears for every candidate already in process, the comparison matrix grows a column, and the rollup recomputes. Nobody opens an editor.

Where the machine stops

Read this one as two lanes. Work moves down the top lane until it reaches something only a person can do, crosses to the bottom lane, and comes back. The crossings are the product.



The guard, in the code

Gate 2 is not a convention anyone can forget. `SIGN_OFF` refuses and explains itself unless all three conditions hold:

<code>unrated === 0</code>	Every competency row carries a rating on the six-point scale. No decimals, as the sheet specifies.
<code>culturalFitment != null</code>	A value from the three-point fitment scale has been chosen.
<code>finalDecision != null</code>	One of the three IAS outcomes is recorded: suitable, on hold, or not suitable.

WHAT THE MACHINE CONTRIBUTES

A suggested rating, the evidence it came from, and a confidence figure. Visible, labelled, and worth nothing until someone agrees with it.

WHAT COUNTS

Only rows a person has confirmed enter `weightedAverage`. An untouched AI draft scores zero, not 4.7.

FLOW 04 – TRACEABILITY

From a sentence to a hiring recommendation

Every number in the analysis screens can be walked backwards to something a candidate actually said. This is one real chain from the sample data, start to finish.

01 – SPOKEN · INT-001 · T+00:53

"First thing I did was stop everyone guessing. I froze changes, put one engineer on data collection only, and got the customer on a 30-minute cadence."
ANANYA SHARMA · CANDIDATE

02 – TAGGED LIVE

competencyId = c1_esc
Technical Escalation Management · weight 0.10

03 – SUGGESTED

rating 6 / 6 · confidence 0.94
carries the quote as its evidence

04 – CONFIRMED

Murugan Iyer accepts 6 · signs the level
interviewerConfirmed = true

05 – SCORECARD AVERAGE · ASM-4101-S_L1

$\Sigma(\text{rating} \times \text{weight}) \div \Sigma(\text{weight}) = 5.20 / 6$
over rated rows only – an unrated row is absent, not zero

06 – CONSOLIDATED · 3 SIGNED LEVELS

5.18 / 6 across L1, L2, L3
unsigned levels contribute nothing and say so

07 – RECOMMENDATION · SHOWN WITH ITS CITATIONS

"Strongest on Technical Escalation Management (6.0) – froze changes and installed a single named owner within the first hour." Rated 6/6 by Murugan Iyer.

FIG 05 Seven hops, no black box. A hiring manager challenging the recommendation at step 07 can be shown the sentence at step 01 in two clicks.

REFERENCE

The record

Six entities carry the whole product. The cardinalities are where the dynamism lives – every 1:N on this diagram is a place the requisition gets to decide N.




FIG 06 The dashed blue edge is figure 03 restated as a relationship: a workflow stage is what causes an assessment to exist.

How state actually moves

One reducer holds everything. A screen dispatches, the reducer returns new entities plus an audit entry, and derived selectors recompute from there. There is no second source of truth to fall out of sync.

ACTION	WHAT IT CHANGES	WRITES AUDIT
SCREEN_DECISION	Candidate status and stage; closes the screening task	Yes, with the AI suggestion alongside

ACTION	WHAT IT CHANGES	WRITES AUDIT
ACCEPT_AI	Copies the draft into the row and marks it confirmed	Yes
SCORE_ROW	One rating or note; recomputes the weighted average	No — captured at sign-off
SIGN_OFF	Locks the scorecard, advances the stage, records the score	Yes, with the count of overrides
ADD_STAGE	Appends a level; downstream generation follows	Yes
APPROVE_FRAMEWORK	Publishes to the repository if weights total 100%	Yes

REFERENCE

Where this plugs into what you already run

Nexora is the interview and decision layer, not a system of record. The planes in figure 01 are drawn the way they are so an existing ATS can own the record without any of this changing.

LIVE TODAY

Invites and reminders by email. Branded PDF export of comparisons and scorecards. Multi-country scheduling in each interview's own timezone.

DESIGNED FOR, NOT WIRED

Calendar sync for panel availability. Native video platform meetings and recording ingestion. ATS requisition and candidate sync.

DELIBERATELY ABSENT

Auto-rejection. Ranking that picks rather than orders. Any path from a model output to a candidate-facing outcome without a signature.

The three questions this model should let anyone answer

Fig 03

“What happens if this role needs six rounds instead of four?”
One array entry. Everything downstream follows.

Fig 04

“Is the AI making hiring decisions?” No, and there is no code path by which it could. Three gates, all blocking.

Fig 05

“Where did this score come from?” Seven hops back to a sentence someone said, with the interviewer's name on it.

Drawn from the implementation, not from intent — every figure here reflects code that runs. Assessment structure derives from the Interview Assessment Sheet v3.1.